

# Cyberbullying Detection on Twitter Data Using Machine Learning Classifiers

Pradip Dhakal

*Statistics and Data Science Department  
University of Central Florida*

**Abstract**—This study compares some of the popular machine learning techniques like Logistic Regression, Multinomial Naive Bayes, K-Nearest Neighbor, and Extreme Gradient Boosting to classify the tweets into three different categories: cyberbullying based on religion, cyberbullying based on ethnicity, or no cyberbullying. First, various data-cleaning approaches are used to clean the tweet data. After the data is clean and ready, the word embedding techniques, such as a bag of words and term frequency-Inverse document frequency, are used to convert the words into mathematical vectors. Finally, the model will be fitted using the combination of the above-mentioned word embedding techniques and machine learning algorithms.

**Keywords**—Logistic Regression, Multinomial Naive Bayes, K-Nearest Neighbor, Extreme Gradient Boosting, Bag of Words, Term Frequency-Inverse Document Frequency

## I. INTRODUCTION

Social media has been popular for quite a long time. People from all around the world are able to communicate with each other, share their knowledge and thoughts, and know what's happening on another side of the world. Some of the popular social platforms are Facebook, Instagram, Twitter, YouTube, and Snapchat. Alongside these advantages, people argue with each other, show aggressive behaviors, leave negative and racist comments, and bully other people on different social platforms. Many people have been victims of these kinds of activities, leading to an adverse psychological impact on the victim's emotions. Many research studies have been proposed to mitigate these kinds of activities. Most researchers formulated this problem as a classification problem. The study by Dinakar et al. [1] performed binary classification to see whether the comments on YouTube could be classified as sensitive or not; they also performed multi-label classification to see what comments belong to what classes. Another study from Chavan and Shylaja [2] performed the binary classification, where they classified the texts as bullying texts and non-bullying texts. Similarly, the study from Dadvar et al. [3] incorporated the user's age and gender to improve the accuracy of cyberbullying detection. These articles motivated me to work in the area of cyberbullying detection.

Our study consists of three different classes: ethnicity-based cyberbullying with the label '1', religion-based cyberbullying with the label '2', and no cyberbullying with the label '0'. Our primary goal is to come up with the best model that classifies the words very accurately. The best model that we obtain from this study can be used in the social application platform to

detect bullying words and warn the users to use alternative words.

## II. DESIGN OVERVIEW

This section explains the methodology and framework used in this study. In fig 1, we can see the flowchart of the overall design. First, the data is cleaned and pre-processed. Second, the cleaned data is used to create the word clouds for each class. Also, the train-test split is performed after the data pre-processing. Next, taking the split data, two different word embedding techniques: Bag of Words (BoW) and Term Frequency-Inverse Document Frequency (TF-IDF), are used to convert the tokens into mathematical vectors. Finally, four different classification algorithms: Logistic Regression (LR), Multinomial Naive Bayes (MNB), K-Nearest Neighbor (KNN), and Extreme Gradient Boosting (XgBoost) are used with each word embedding technique to perform the multi-label classification. The results are then compared using the classification accuracy and F1 score metrics.

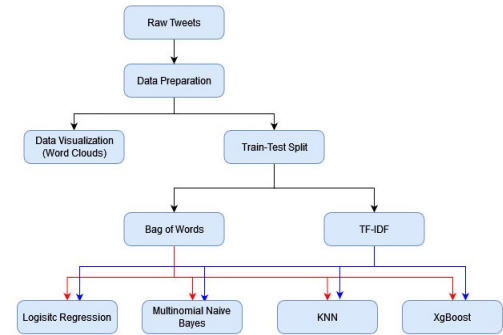


Fig. 1. Design Overview

## III. DATA

### A. Data Information

The dataset used in this paper is available here to the public, and it was contributed by J. Wang et al. [4]. Among the available various classes, I am using the text data sets that belong to ethnicity, religion, and not cyberbullying classes.

## 1. Ethnicity-Based Cyberbullying Class



### 3.Non Cyberbullying Class



Logistic Regression		
Word Embedding	Accuracy	F1 Score
BoW	0.96	0.96
TF-IDF	0.96	0.96

TABLE I  
LOGISTIC REGRESSION RESULTS

From the table I, both BoW and TF-IDF gave the same classification accuracy and F1 score with the logistic regression model.

Confusion Matrices:

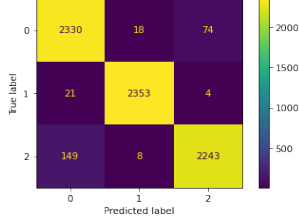


Fig. 5. Logistic Regression with Bag of Words

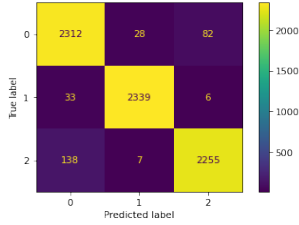


Fig. 6. Logistic Regression with TF-IDF

## 2. Multinomial Naive Bayes

The multinomial Naive Bayes algorithm is the variation of the traditional naive Bayes algorithm that is employed in text classification. The naive Bayes algorithm assumes conditional independence between each pair of features given the class variable. The classification rule for naive Bayes is given below.

$$P(y|x_1, x_2, \dots, x_n) \propto P(y) \prod_{i=1}^n P(x_i|y)$$

$$\hat{y} = \underset{y}{argmax} P(y) \prod_{i=1}^n P(x_i|y)$$

The multinomial naive Bayes model was fitted with both BoW and TF-IDF on the training dataset. On predicting the results on the test dataset, the following results were obtained.

Multinomial Naive Bayes		
Word Embedding	Accuracy	F1 Score
BoW	0.88	0.87
TF-IDF	0.86	0.85

TABLE II  
MULTINOMIAL NAIVE BAYES RESULTS

From the table II, the BoW performed slightly better than TF-IDF with the multinomial naive Bayes algorithm.

Confusion Matrices:

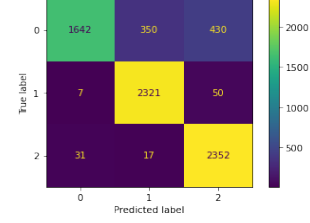


Fig. 7. Multinomial Naive Bayes with Bag of Words

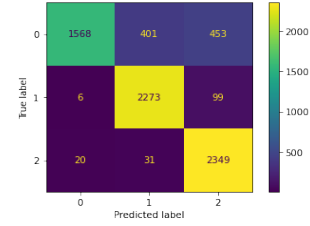


Fig. 8. Multinomial Naive Bayes with TF-IDF

## 3. K-Nearest Neighbor

This is another classification algorithm that finds the k nearest neighbors of a new point and labels the new point to the class which has a high number of points out of the k neighbors. The predicted value is given by the following equation.

$$\hat{y} = \frac{1}{K} \sum_{x_i \in N_k(x)} y_i \quad (3)$$

where  $N_k(x)$  indicates k samples from the training data that are closest to x.

Using the equation 3, we can predict the class of each observation. The Minkowski distance is used to calculate the distance between the points to get the closest neighbors.

Before fitting the model, the optimal value for k needs to be calculated using both BoW and TF-IDF data. I used a 5-fold cross-validation method and obtain the optimal value for k to be 1 for both of them. The KNN model is now fitted with both BoW and TF-IDF on the training dataset. On predicting the results on the test dataset, the following results were obtained.

K-Nearest Neighbor			
Word Embedding	Best K	Accuracy	F1 Score
BoW	1	0.86	0.86
TF-IDF	1	0.46	0.40

TABLE III  
K-NEAREST NEIGHBOR RESULTS

From the table III, the BoW performed way better than the TF-IDF with the K-Nearest Neighbor algorithm.

Confusion Matrices:

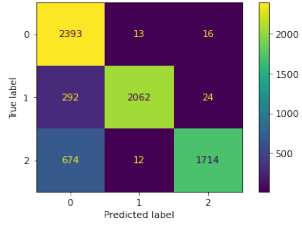


Fig. 9. K-Nearest Neighbor with Bag of Words

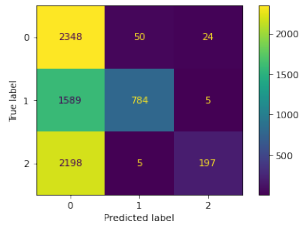


Fig. 10. K-Nearest Neighbor with TF-IDF

#### 4. Extreme Gradient Boosting

Extreme gradient boosting is an improved version of Gradient Boosting Machines that creates more stable models by drastically lowering the likelihood of overfitting. Extreme gradient boosting is an ensemble learning technique that is basically built on the concept of a random forest algorithm. Ensemble learning is a method where multiple learners are employed to obtain better prediction accuracy; the model learns from previous mistakes at each iteration and reduces the loss.

Let  $F = f_1, f_2, \dots, f_n$  be the set of base learners, then the final prediction is given by

$$\hat{y}_i = \sum_{j=1}^n f_j(x_i)$$

A base learner is fitted to the negative gradient of the loss function with respect to the value from the previous iteration in order to get  $f_j(x_i)$  at each iteration.

The xgboost model was fitted with both BoW and TF-IDF on the training dataset using the default parameters. On predicting the results on the test dataset, the following results were obtained.

XgBoost		
Word Embedding	Accuracy	F1 Score
BoW	0.97	0.97
TF-IDF	0.97	0.97

TABLE IV  
XGBOOST RESULTS

Confusion Matrices:

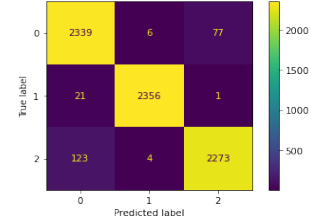


Fig. 11. Xgboost with Bag of Words

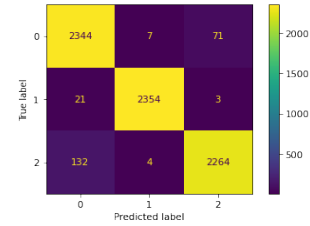


Fig. 12. Xgboost with TF-IDF

Next, a couple of hyperparameters: max depth and learning rate, were tuned to see if we could further improve the accuracy results.

Tuned Hyperparameters		
Word Embedding	max depth	learning rate
BoW	4	0.5
TF-IDF	6 (same as default)	0.3 (same as default)

TABLE V  
XGBOOST HYPERPARAMETER TUNING

Table V lists the value of tuned max depth and learning rate. The tuned hyperparameter values were the same as the default values for TF-IDF; however, the tuned hyperparameter values for BoW differed. The xgboost model was again fitted with BoW using the tuned hyperparameter values, and the following result was obtained.

XgBoost with Tuned Hyperparameters		
Word Embedding	Accuracy	F1 Score
BoW	0.97	0.97
TF-IDF	0.97	0.97

TABLE VI  
XGBOOST WITH TUNED HYPERPARAMETERS RESULTS

The accuracy and F1 score in table IV and table VI are the same. The tuning of the hyperparameters did not increase the

accuracy and F1 score of the xgboost model. Also, both BoW and TF-IDF performed the same with the xgboost algorithm.

Confusion Matrix:

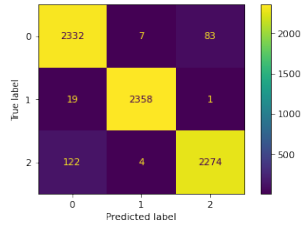


Fig. 13. Xgboost using Tuned Hyperparameters with Bag of Words

## V. CONCLUSION

Comparing all four models, the highest accuracy and f1 score of 97% were obtained from the xgboost model with both BoW and TF-IDF. The second highest accuracy and f1 score of 96% were obtained from the logistic regression with both BoW and TF-IDF. The results are summarized in the table given below.

All Results					
Word Embedding	Metrics	LR	MNB	KNN	XgBoost
BoW	Accuracy	<b>0.96</b>	0.88	0.86	<b>0.97</b>
	F1 Score	<b>0.96</b>	0.87	0.86	<b>0.97</b>
TF-IDF	Accuracy	<b>0.96</b>	0.86	0.46	<b>0.97</b>
	F1 Score	<b>0.96</b>	0.85	0.40	<b>0.97</b>

TABLE VII  
FINAL COMPARISON

The accuracy and f1 score is not much different for the xgboost and the logistic regression model. For this case study, accuracy matters more than interpretability, so even though the xgboost model is more complex than the logistic regression model, the xgboost model is preferred.

## REFERENCES

- [1] K. Dinakar, R. Reichart, and H. Lieberman, "Modeling the detection of textual cyberbullying," In Proceedings of the Social Mobile Web. Citeseer, 2011.
- [2] V. S. Chavan and S. Shylaja, "Machine learning approach for detection of cyber-aggressive comments by peers on social media network," in 2015 International Conference on Advances in Computing, Communications and Informatics (ICACCI). IEEE, 2015, pp. 2354–2358.
- [3] M. Dadvar, F. d. Jong, R. Ordelman, and D. Trieschnigg, "Improved cyberbullying detection using gender information," in Proceedings of the Twelfth Dutch-Belgian Information Retrieval Workshop (DIR 2012). University of Ghent, 2012.
- [4] J. Wang, K. Fu, C.T. Lu, "SOSNet: A Graph Convolutional Network Approach to Fine-Grained Cyberbullying Detection," Proceedings of the 2020 IEEE International Conference on Big Data (IEEE BigData 2020), pp. 1699-1708, December 10-13, 2020.

## APPENDIX

All the python code for this work is given below:

```
# Importing all the required packages

import pandas as pd
import numpy as np
import re
import matplotlib.pyplot as plt
import nltk
from nltk.stem import WordNetLemmatizer
from sklearn.preprocessing import LabelEncoder
import string
import re
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
from sklearn.linear_model import LogisticRegression
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import CountVectorizer
from sklearn import metrics
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
from sklearn.naive_bayes import MultinomialNB
from sklearn.neighbors import KNeighborsClassifier
import xgboost as xgb
from wordcloud import WordCloud
from sklearn.model_selection import GridSearchCV

nltk.download('stopwords')
nltk.download('punkt')
nltk.download('wordnet')
nltk.download('omw-1.4')
# Loading Data

def openTextasList(filename):
```



```
with open(filename, encoding="utf8") as file_in:
```

```
    lines = []
```

```
    for line in file_in:
```

```
        # remove whitespace characters like `\\n` at the end of each line
```

```
        line=line.strip()
```

```
        lines.append(line)
```

```
    return(lines)
```

```
ethnicity=openTextasList("../DataMiningIFinalProject/8000ethnicity.txt")
```

```
religion=openTextasList("../DataMiningIFinalProject/8000religion.txt")
```

```
notcb=openTextasList("../DataMiningIFinalProject/8000notcb.txt")
```

```
alldoc=notcb + ethnicity + religion
```

```
print("There are %d ethnicity tweets:\\n %s " % (len(ethnicity),ethnicity[0:5]))
```

```
print("There are %d religion tweets:\\n %s " % (len(religion),religion[0:5]))
```

```
print("There are %d notcb tweets:\\n %s " % (len(notcb),notcb[0:5]))
```

```
alldoc[0:2]
```

```
# define training labels
```

```
class_label = np.array(["notcb"for _ in range(8000)] +
```

```
                        ["ethencity"for _ in range(8000)] +
```

```
                        ["religion"for _ in range(8000)])
```

```
class_label.shape
```

```
# Construct a dataframe
```

```
lst = alldoc
```

```
# list of int
```

```
lst2 = class_label
```

```
# zipping both lists with columns specified
```

```
df = pd.DataFrame(list(zip(lst, lst2)),
```

```
                    columns=["Tweets", 'Labels'])
```



```
df
```

```
df['Labels'].value_counts()
```

```
# Data Cleaning
```

```
np.sum(df.isnull())
```

```
# storing both Tweets and Labels in lists
```

```
tweets, labels = list(df['Tweets']), list(df['Labels'])
```

```
# labels
```

```
labelencoder = LabelEncoder()
```

```
df['LabelsEncoded'] = labelencoder.fit_transform(df['Labels'])
```

```
#Changing notcb to 0
```

```
df.LabelsEncoded = df.LabelsEncoded.replace([0,1,2], [1,0,2])
```

```
df[['Labels', 'LabelsEncoded']].value_counts()
```

```
# converting tweets to lower case
```

```
df['Tweets'] = df['Tweets'].str.lower()
```

```
df.head()
```

```
# removing stopwords
```

```
def RemoveStopWords(input_text):
```

```
    StopWordsList = stopwords.words('english')
```

```
    # Words that might indicate some sentiments are assigned to
```

```
    # WhiteList and are not removed
```

```
    WhiteList = ["n't", "not", "no"]
```

```
    words = input_text.split()
```

```
    CleanWords = [word for word in words if (word not in StopWordsList or word in WhiteList) and len(word) > 1]
```

```
    return " ".join(CleanWords)
```

```
df.Tweets = df["Tweets"].apply(RemoveStopWords)
```

```
df.Tweets.head()
```

```
# removing URLs
```

```
def RemoveURLs(text):
```

```
    return re.sub(r"((www.[^s]+)|(http\S+))", "", text)
```

```
df["Tweets"] = df["Tweets"].apply(lambda x : RemoveURLs(x))
```

```
df.Tweets.head()
```

```
# removing mentions
```

```
def MentionsRemover(input_text):
```

```
    return re.sub(r'@\w+', "", input_text)
```

```
df.Tweets = df["Tweets"].apply(MentionsRemover)
```

```
df.Tweets.head()
```

```
# removing numeric data
```

```
def RemoveNumeric(text):
```

```
    return re.sub('[0-9]+', "", text)
```

```
df["Tweets"] = df["Tweets"].apply(lambda x: RemoveNumeric(x))
```

```
df.Tweets.head()
```

```
# removing punctuations
```

```
Punctuations = string.punctuation
```

```
print(Punctuations)
```

```
def RemovePunctuations(text):
```

```
    translator = str.maketrans("", "", Punctuations)
```

```
    return text.translate(translator)
```

```
df['Tweets'] = df['Tweets'].apply(lambda x : RemovePunctuations(x))
df.Tweets.head()
```

```
# removing emojis
```

```
def RemoveEmoji(text):
```

```
    EmojiPattern = re.compile(pattern = "["
                                u"\U0001F600-\U0001F64F" # emoticons
                                u"\U0001F300-\U0001F5FF" # symbols & pictographs
                                u"\U0001F680-\U0001F6FF" # transport & map symbols
                                u"\U0001F1E0-\U0001F1FF" # flags (iOS)
                                "]" + , flags = re.UNICODE)
    return EmojiPattern.sub(r'',text)
```

```
df['Tweets'] = df['Tweets'].apply(lambda x : RemoveEmoji(x))
df.Tweets.head()
```

```
# Tokenization of tweets
```

```
df.Tweets = df.Tweets.tolist()
```

```
TokenizeText = [word_tokenize(i) for i in df.Tweets]
```

```
# for i in TokenizeText:
```

```
#     print(i)
```

```
df.Tweets = TokenizeText
```

```
print(df.Tweets.head())
```

```
# Lemmatization
```

```
lemmatizer = WordNetLemmatizer()
```

```
def Lemmatization(text):
```

```
    text = [lemmatizer.lemmatize(word) for word in text]
```

```
    return text
```

```
df['Tweets'] = df['Tweets'].apply(lambda x: Lemmatization(x))
print(df['Tweets'].head())
```

```
df
```

```
# Joining all words with spaces
```

```
df['Tweets'] = df['Tweets'].apply(lambda x : " ".join(x))
df
```

```
# Word Clouds
```

```
# NOT cyberbullying tweets
```

```
NotCbDf = df.loc[df['LabelsEncoded'] == 0]
```

```
# Converting all tweets into a single list and then to single string
```

```
NotCbDfTweets = NotCbDf.Tweets.tolist()
```

```
NotCbDfTweets = " ".join(NotCbDfTweets)
```

```
#pip install wordcloud
```

```
# Word Cloud for NOT cyberbullying tweets
```

```
wordcloud = WordCloud(max_words=50).generate(NotCbDfTweets)
```

```
plt.figure(figsize=(12, 8))
```

```
plt.imshow(wordcloud, interpolation='bilinear')
```

```
plt.axis("off")
```

```
plt.show()
```

```
# Ethnicity tweets
```

```
EthnicityDf = df.loc[df['LabelsEncoded'] == 1]
```

```
# Converting all tweets into a single list and then to single string
```

```
EthnicityDfTweets = EthnicityDf.Tweets.tolist()
```

```
EthnicityDfTweets = " ".join(EthnicityDfTweets)
```

```
# Word Cloud for ethnicity tweets
```

```
wordcloud = WordCloud(max_words=50).generate(EthnicityDfTweets)
```

```
plt.figure(figsize=(12, 8))
```

```
plt.imshow(wordcloud, interpolation='bilinear')
```

```
plt.axis("off")
```

```
plt.show()
```

```
# Religion tweets
```

```
ReligionDf = df.loc[df['LabelsEncoded'] == 2]
```

```
# Converting all tweets into a single list and then to single string
```

```
ReligionDfTweets = ReligionDf.Tweets.tolist()
```

```
ReligionDfTweets = " ".join(ReligionDfTweets)
```

```
# Word Cloud for religion tweets
```

```
wordcloud = WordCloud(max_words=50).generate(ReligionDfTweets)
```

```
plt.figure(figsize=(12, 8))
```

```
plt.imshow(wordcloud, interpolation='bilinear')
```

```
plt.axis("off")
```

```
plt.show()
```

```
# Splitting the data into train-test
```

```
# Splitting the data
```

```
X, y = df['Tweets'], df['LabelsEncoded']
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size= 0.3, random_state= 111)
```

```
print(X_train.shape)
```

```
print(X_test.shape)
```

```
# Bag of Words and TF-IDF
```

```

# Bag of Words

BoW = CountVectorizer(ngram_range= (1,1))

# Train data

BoW_X_train = BoW.fit_transform(X_train)

print(BoW_X_train.toarray())

print(BoW_X_train.toarray().shape)

# Test data

BoW_X_test = BoW.transform(X_test)

print(BoW_X_test.toarray())

print(BoW_X_test.toarray().shape)


#Check

BoW_X_train.toarray()[100][21900:21950]


# TF-IDF

TF_IDF = TfidfVectorizer(ngram_range=(1,1), max_features= 200000)

#Train Data

TF_IDF_X_train = TF_IDF.fit_transform(X_train)

print(TF_IDF_X_train.toarray())

print(TF_IDF_X_train.toarray().shape)

# Test Data

TF_IDF_X_test = TF_IDF.transform(X_test)

print(TF_IDF_X_test.toarray())

print(TF_IDF_X_test.toarray().shape)


#Check

TF_IDF_X_train.toarray()[100][21900:21950]


# Modeling


# Logistic Regression

```

```

# Function for logistic regression to compare Bag of Words and TF-IDF
def LogisticRegressionFunction(X_train, X_test, y_train, y_test, description):
    LogitClassifier = LogisticRegression(solver='lbfgs', multi_class='multinomial', random_state=111, n_jobs=-1)
    LogitClassifier.fit(X_train, y_train)
    y_prediction = LogitClassifier.predict(X_test)
    ConfMat = confusion_matrix(y_test, y_prediction)
    display = ConfusionMatrixDisplay(confusion_matrix= ConfMat)
    display.plot()
    plt.show()
    Accuracy = metrics.accuracy_score(y_test, y_prediction)
    F1 = metrics.f1_score(y_test, y_prediction, average='weighted')
    print('Accuracy for ', description, ' is: {:.2f}'.format(Accuracy))
    print('F1 score for ', description, ' is: {:.2f}'.format(F1))

# Logistic regression with Bag of Words
LogisticRegressionFunction(BoW_X_train, BoW_X_test, y_train, y_test, 'Logistic Regression with Bag of Words')

# Logistic regression with TF-IDF
LogisticRegressionFunction(TF_IDF_X_train, TF_IDF_X_test, y_train, y_test, 'Logistic regression with TF-IDF')

# Multinomial Naive Bayes

# Function for multinomial naive bayes to compare Bag of Words and TF-IDF
def MultinomialNaiveBayes(X_train, X_test, y_train, y_test, description):
    MultiNaiveBayesClassifier = MultinomialNB()
    MultiNaiveBayesClassifier.fit(X_train, y_train)
    y_prediction = MultiNaiveBayesClassifier.predict(X_test)
    ConfMat = confusion_matrix(y_test, y_prediction)
    display = ConfusionMatrixDisplay(confusion_matrix= ConfMat)
    display.plot()
    plt.show()

```



```

Accuracy = metrics.accuracy_score(y_test, y_prediction)

F1 = metrics.f1_score(y_test, y_prediction, average='weighted')

print('Accuracy for ', description, ' is: {:.2f}'.format(Accuracy))

print('F1 score for ', description, ' is: {:.2f}'.format(F1))


# Multinomial Naive Bayes with Bag of Words

MultinomialNaiveBayes(BoW_X_train, BoW_X_test, y_train, y_test, 'Multinomial Naive Bayes with Bag of Words')


# Multinomial Naive Bayes with TF-IDF

MultinomialNaiveBayes(TF_IDF_X_train, TF_IDF_X_test, y_train, y_test, 'Multinomial Naive Bayes with TF-IDF')


# K-Nearest Neighbor


# First, lets find the best value of k.


# define grid parameters

grid_params = { 'n_neighbors': list(range(1,26))}

# grid search

GS = GridSearchCV(KNeighborsClassifier(), grid_params, verbose = 1, cv=5, n_jobs = -1)


# fit the model for Bag of Words

GridResult = GS.fit(BoW_X_train, y_train)

# hyperparameters with the best score

GridResult.best_params_


# fit the model for Bag of Words

GridResult = GS.fit(TF_IDF_X_train, y_train)

# hyperparameters with the best score

GridResult.best_params_

```

```
# Function for knn to compare Bag of Words and TF-IDF
def KnnFunction(X_train, X_test, y_train, y_test, n, description):
```

```
    KnnClassifier = KNeighborsClassifier(n_neighbors= n)
    KnnClassifier.fit(X_train, y_train)
    y_prediction = KnnClassifier.predict(X_test)
    ConfMat = confusion_matrix(y_test, y_prediction)
    display = ConfusionMatrixDisplay(confusion_matrix= ConfMat)
    display.plot()
    plt.show()
    Accuracy = metrics.accuracy_score(y_test, y_prediction)
    F1 = metrics.f1_score(y_test, y_prediction, average='weighted')
    print('Accuracy for ', description, ' is: {:.2f}'.format(Accuracy))
    print('F1 score for ', description, ' is: {:.2f}'.format(F1))
```

```
# KNN with Bag of Words
```

```
KnnFunction(BoW_X_train, BoW_X_test, y_train, y_test, 1, 'KNN with Bag of Words')
```

```
# KNN with TF-IDF
```

```
KnnFunction(TF_IDF_X_train, TF_IDF_X_test, y_train, y_test, 1, 'KNN with TF-IDF')
```

```
# Extreme Gradient Boosting
```

```
# Function for xgboost to compare Bag of Words and TF-IDF
```

```
def XgBoostFunction(X_train, X_test, y_train, y_test, learning_rate, max_depth, description):
    XgBoostClassifier = xgb.XGBClassifier(objective = 'multi:softmax',
                                           learning_rate = learning_rate, max_depth = max_depth, seed = 111)
    XgBoostClassifier.fit(X_train, y_train)
    y_prediction = XgBoostClassifier.predict(X_test)
    ConfMat = confusion_matrix(y_test, y_prediction)
    display = ConfusionMatrixDisplay(confusion_matrix= ConfMat)
    display.plot()
    plt.show()
```

```

Accuracy = metrics.accuracy_score(y_test, y_prediction)

F1 = metrics.f1_score(y_test, y_prediction, average='weighted')

print('Accuracy for ', description, ' is: {:.2f}'.format(Accuracy))

print('F1 score for ', description, ' is: {:.2f}'.format(F1))


# The default value for learning_rate is 0.3 and max_depth is 6.


# Xgboost with Bag of Words (using default parameters)

XgBoostFunction(BoW_X_train, BoW_X_test, y_train, y_test, 0.3, 6, 'Xgboost with Bag of Words')


# Xgboost with TF-IDF (Using default parameters)

XgBoostFunction(TF_IDF_X_train, TF_IDF_X_test, y_train, y_test, 0.3, 6, 'Xgboost with TF-IDF')


# Let's tune a couple of parameters: learning_rate and max_depth.


# define parameters

params = { 'max_depth': [3, 4, 5, 6, 7],
           'learning_rate': [0.01, 0.1, 0.3, 0.5, 0.7] }

# grid search

GrdSrch = GridSearchCV(estimator= xgb.XGBClassifier(objective = 'multi:softmax', seed = 111),
                       param_grid= params,
                       scoring='accuracy',
                       verbose=1)


# fit the model for Bag of Words

GrdSrchResult = GrdSrch.fit(BoW_X_train, y_train)

# hyperparameters with the best score

GrdSrchResult.best_params_


# fit the model for TF-IDF

GrdSrchResult = GrdSrch.fit(TF_IDF_X_train, y_train)

# hyperparameters with the best score

```

```
GrdSrchResult.best_params_
```

```
# Xgboost with Bag of Words (using tuned parameters)
```

```
XgBoostFunction(BoW_X_train, BoW_X_test, y_train, y_test, 0.5, 4, 'Xgboost with Bag of Words')
```

```
# For TF-IDF, we got default values as the best ones.
```